Alternative Routes in Road Networks

ITTAI ABRAHAM, DANIEL DELLING, ANDREW V. GOLDBERG
and RENATO F. WERNECK, Microsoft Research Silicon Valley

We study the problem of finding good alternative routes in road networks. We look for routes that are substantially different from the shortest path, have small stretch, and are locally optimal. We formally define the problem of finding alternative routes with a single via vertex, develop efficient algorithms for it, and evaluate them experimentally. Our algorithms are efficient enough for practical use and compare favorably with previous methods in both speed and solution quality.

Categories and Subject Descriptors: G.2.2 [**Graph Theory**]: Graph algorithms

General Terms: Algorithms, Experimentation, Measurement, Performance

Additional Key Words and Phrases: shortest paths, route planning, alternative paths, speedup techniques

1. INTRODUCTION

We use web-based and autonomous navigation systems in everyday life. These systems produce driving directions by computing a shortest path (or an approximation) with respect to a length function based on various measures, such as distance and travel time. However, optimal paths do not necessarily match the personal preferences of individual users. These preferences may be based on better local knowledge, a bias for or against a certain route segment, or other factors. One can deal with this issue by presenting a small number of alternative paths and hoping one of them will satisfy the user. The goal of an algorithm is to offer alternative paths often, and for these alternatives to look reasonable.

Recent research on route planning focused on computing only a single (shortest) path between two given vertices (see [Delling et al. 2009] for an overview). Much less work has been done on finding multiple routes. Some commercial products (by companies such as Google and TomTom) suggest alternative routes using proprietary algorithms. Among published results, a natural approach is to use k-shortest path algorithms [Eppstein 1994], but this is impractical because a reasonable alternative in a road network is probably not among the first few thousand paths. Another approach is to use multi-criteria optimization [Hansen 1979; Martins 1984], in which two or more length functions are optimized at once, and several combinations are returned. Efficient algorithms have been recently presented by Delling and Wagner [2009] and Geisberger et al. [2010]. Our focus is on computing reasonable alternatives with a single length function. In this context, the best published results we are aware of are produced by the *choice routing* algorithm [Cambridge Vehicle Information Technology Ltd. 2005], which we discuss in Section 4. Although it produces reasonable paths, it is not fully documented and is not fast enough for continental-sized networks.

Authors' addresses: Ittai Abraham, Daniel Delling, Andrew V. Goldberg, and Renato F. Werneck, Microsoft Research Silicon Valley, Mountain View, 94043, USA; email: {*ittaia, dadellin, goldberg, renatow*}*@microsoft.com*
Permission to make digital or hard copies of part or all of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies show this notice on the first page or initial screen of a display along with the full citation. Copyrights for components of this work owned by others than ACM must be honored. Abstracting with credit is permitted. To copy otherwise, to republish, to post on servers, to redistribute to lists, or to use any component of this work in other works requires prior specific permission and/or a fee. Permissions may be requested from Publications Dept., ACM, Inc., 2 Penn Plaza, Suite 701, New York, NY 10121-0701 USA, fax +1 (212) 869-0481, or permissions@acm.org.
© ???? ACM 1084-6654/????/03-ARTZ $15.00
DOI 10.1145/0000000.0000000 http://doi.acm.org/10.1145/0000000.0000000

Fig. 1. Example query output by our most efficient algorithm. The source has label (1) and the target is labeled (2). The quickest path is drawn red (4), the first alternative blue (6), and the second alternative green (8). In addition, (6) and (8) label the via vertices of the alternatives. (Visualization by Bing Maps.)

In this work, we study the problem of finding "reasonable" alternative routes in road networks. We start by defining in Section 3 a natural class of *admissible* alternative paths. Obviously, an alternative route must be substantially different from the optimal path, and must not be much longer. But this is not enough: alternatives must feel natural to the user, with no unnecessary detours. (Formally, they must be *locally optimal*: every subpath up to a certain length must be a shortest path.) Even with these restrictions, the number of admissible paths may be exponential, and computing the best one is hard. For efficiency, we focus on a more limited (yet useful) subset. Given an origin s and a target t, we restrict ourselves to *single via paths*, which are alternative routes obtained by concatenating the shortest paths from s to v and from v to t, for some vertex v. Section 4 discusses how the best such path can be computed in polynomial time using the bidirectional version of Dijkstra's algorithm (BD).

In practice, however, just polynomial time is not enough—we need sublinear algorithms. Modern algorithms for computing (optimal) shortest paths in road networks, such as those based on contraction hierarchies [Geisberger et al. 2008] and reach [Goldberg et al. 2009], are often based on pruning Dijkstra's search. After practical preprocessing steps, they need to visit just a few hundred vertices to answer queries on continental-sized graphs with tens of millions of vertices—orders of magnitude faster than BD. In fact, as shown by Abraham, Fiat et al. [2010b], their performance is sublinear on graphs with small highway dimension, such as road networks. Section 5 shows how to apply these speedup techniques to the problem of finding alternative routes, and Section 6 proposes additional measures to make the resulting algorithms truly practical.

Finally, Section 7 evaluates various algorithms experimentally according to several metrics, including path quality and running times. We show that finding a good alternative path takes only six times as much as computing the shortest path (with a pruning algorithm). Moreover, our pruning methods have similar success rates to a variant of choice routing, but are orders of magnitude faster. Figure 1 gives examples of alternative routes computed by our approach. See also Figure 6, in the Appendix.

Summarizing, our contributions are twofold. First, we establish a rigorous theoretical foundation for alternative paths, laying the ground for a systematic study of the

problem. Second, we present efficient algorithms (in theory and in practice) for finding such routes.

This paper is based on an extended abstract presented at the 9th International Symposium on Experimental Algorithms [2010a]. In addition to our original contributions, here we introduce a new relaxation of the algorithm based on contraction hierarchies, provide more experiments, and compare our methods to a natural algorithm based on iteratively penalizing shortest paths until a satisfactory alternative is found.

2. DEFINITIONS AND BACKGROUND

Let $G = (V, E)$ be a directed graph with nonnegative, integral lengths on edges, with $|V| = n$ and $|E| = m$. Given any path P in G, let $|P|$ be its number of edges and $\ell(P)$ be the sum of the lengths of its edges. By extension, $\ell(P \cap Q)$ is the sum of the lengths of the edges shared by paths P and Q, and $\ell(P \setminus Q)$ is $\ell(P) - \ell(P \cap Q)$. Given two vertices, s and t, the *point-to-point shortest path problem* (P2P) is that of finding the shortest path (denoted by $Opt(s, t = Opt)$) from s to t. To simplify the exposition, throughout this paper we assume Opt is unique; this can be enforced by breaking ties consistently or by adding small perturbations to the graph.

Dijkstra's algorithm [Dijkstra 1959] computes $\text{dist}(s, t)$ (the distance from s to t in G) by scanning vertices in increasing order from s. Bidirectional Dijkstra (BD) runs a second search from t as well, and stops when both search spaces meet [Dantzig 1962].

The *reach* of v, denoted by $r(v)$, is defined as the maximum, over all shortest u–w paths containing v, of $\min\{\text{dist}(u, v), \text{dist}(v, w)\}$. During an s–t search, BD can be pruned at all vertices v for which both $\text{dist}(s, v) > r(v)$ and $\text{dist}(v, t) > r(v)$ hold [Gutman 2004]. The insertion of *shortcuts* (new edges representing shortest paths in the original graph) may decrease the reach of some original vertices, thus significantly improving the efficiency of this approach [Goldberg et al. 2009]. Queries in the the resulting algorithm (called RE) are three orders of magnitude faster than plain BD on continental-sized road networks.

An even more efficient algorithm (by another order of magnitude) is *contraction hierarchies* (CH) [Geisberger et al. 2008]. During preprocessing, it sorts all vertices by importance (heuristically), then shortcuts them in this order. (To *shortcut* a vertex, we remove it from the graph and add as few new edges as necessary to preserve distances.) The output of the preprocessing procedure consists of the original graph augmented with shortcuts, together with the order itself. A query only follows an edge (u, v) if v is more important than u.

Both RE and CH are correct on any graph with nonnegative lengths, but are effective mostly on road networks. Recently, Abraham, Fiat et al. [2010b] provided a theoretical justification for this empirical observation. They prove that these algorithms work well on graphs with low *highway dimension*, which includes road networks. Roughly speaking, a graph has low highway dimension if, for every $r > 0$ and every ball B of radius $2r$, there exists a small set of vertices that hits all shortest paths of length at least r in B. The intuition behind this is that long shortest paths are covered by few vertices.

3. ADMISSIBLE ALTERNATIVE PATHS

In this paper, we are interested in finding an alternative path P between s and t. By definition, such a path must be significantly different from Opt: the total length of the edges they share must be a small fraction of $\ell(Opt)$.

This is not enough, however. The path must also be *reasonable*, with no unnecessary detours. While driving along it, every local decision must make sense. To formalize this notion, we require paths to be *locally optimal*. A first condition for a path P to be T *locally optimal* (T-LO) is that every subpath P' of P with $\ell(P') \leq T$ must be a

shortest path. This would be enough if P were continuous, but for actual (discrete) paths in graphs we must "round up" with a second condition. If P' is a subpath of P with $\ell(P') > T$ and $\ell(P'') < T$ (where P'' is the path obtained by removing the endpoints of P'), then P' must be a shortest path. Note that a path that is not locally optimal includes a local detour, which in general is not desirable. (Users who need a detour could specify it separately.)

Although local optimality is necessary for a path to be reasonable, it is arguably not sufficient (see Figure 2). We also require that every subpath of an alternative paths has good stretch. Formally, we say that a path P has $(1+\epsilon)$ *uniformly bounded stretch* ($(1+\epsilon)$-UBS) if for every subpath P' of P (including P itself) with endpoints s', t' we have $\ell(P') \leq (1+\epsilon)\ell(Opt(s', t'))$ (the path P' from s' to t' has stretch $1+\epsilon$ relative to the shortest path between s' and t').

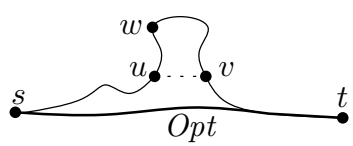

Fig. 2. Rationale for UBS. The alternative through w is a concatenation of two shortest paths, s–w and w–t. Although it has high local optimality, it looks unnatural because there is a much shorter path between u and v.

Given these definitions, we are now ready to define formally the class of paths we are looking for. We need three parameters: $0 < \alpha < 1$, $\epsilon \geq 0$, and $0 \leq \gamma \leq 1$. Given a shortest path Opt between s and t, we say that an s–t path P is an *admissible alternative* if it satisfies the following conditions:

(1) $\ell(Opt \cap P) \leq \gamma \cdot \ell(Opt)$ (limited sharing);
(2) P is T-locally optimal for $T = \alpha \cdot \ell(Opt)$ (local optimality);
(3) P is $(1+\epsilon)$-UBS (uniformly bounded stretch).

There may be zero, one, or multiple admissible alternatives, depending on the input and the choice of parameters. If there are multiple alternatives, we can sort them according to some objective function $f(\cdot)$, which may depend on any number of parameters (possibly including α, ϵ, and γ). In our experiments, we prefer admissible paths with low stretch, low sharing, and high local optimality, as explained in Section 6. Other objective functions could be used as well.

Note that our definitions can be easily extended to report multiple alternative paths. We just have to ensure that the ith alternative is sufficiently different from the union of Opt and all $i-1$ previous alternatives. The stretch and local optimality conditions do not change, as they do not depend on other paths.

4. SINGLE VIA PATHS

Even with the restrictions we impose on admissible paths, they may still be too numerous, making it hard to find the best one efficiently. This section defines a subclass of admissible paths (called *single via paths*) that is more amenable to theoretical analysis and practical implementation. Given any vertex v, the *via path through* v, P_v, is the concatenation of two shortest paths, s–v and v–t. (Recall that we are looking for s–t paths.) We look for via paths that are admissible. As we will see, these can be found efficiently and work well in practice.

Note that single via paths have interesting properties. Among all s–t paths through v (for any v), P_v is the shortest, i.e., it has the lowest stretch. Moreover, being a con-

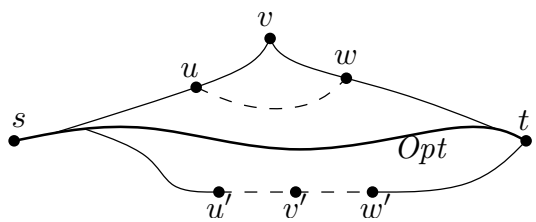

Fig. 3. Example for two T-tests. The T-test for v fails because the shortest path from u to w, indicated as a dashed spline, does not contain v. The test for v' succeeds because the shortest path from u' to w' contains v'.

catenation of two shortest paths, the local optimality of P_v can only be violated around v. In this sense, via paths are close to being admissible.

All $n-2$ via paths can be implicitly generated with a single run of BD, as long as we let each search visit the entire graph. This takes $O(m + n \log n)$ time. Not all via paths found must be admissible, however. For each via path P_v, we must check whether the three admissibility conditions are obeyed.

The easiest condition to check is sharing. For any vertex v visited by the searches, let $\sigma_f(v)$ be the sharing amount in the forward direction (i.e., how much s–v shares with Opt, which is known). Set $\sigma_f(s) \leftarrow 0$ and, for each vertex v (in forward scanning order), set $\sigma_f(v)$ to $\sigma_f(p_f(v)) + \ell(p_f(v), v)$ if $v \in Opt$ or to $\sigma_f(p_f(v))$ otherwise (here p_f denotes the parent in the forward search). Computing $\sigma_r(v)$, the sharing in the reverse direction, is similar. The total sharing amount $\sigma(v) = \ell(Opt \cap P_v)$ is $\sigma_f(v) + \sigma_r(v)$. Note that each vertex in the search space can be processed in constant time, which means this procedure can compute $\sigma(v)$ for all vertices v in $O(n)$ total time (excluding the time to run BD).

In contrast, stretch and local optimality are much harder to evaluate, requiring quadratically many shortest path queries (on various pairs of vertices). Ideally, we would like to verify whether a path P is locally optimal (or is $(1+\epsilon)$-UBS) in time proportional to $|P|$ and a few shortest-path queries. We do not know how to do this. Instead, we present alternative tests that are efficient, have good approximation guarantees, and work well in practice.

For local optimality, there is a quick 2-approximation. Take a via path P_v and a parameter T. Let P_1 and P_2 be the s–v and v–t subpaths of P_v, respectively. Among all vertices in P_1 that are at least T away from v, let x be the closest to v (and let $x = s$ if $\ell(P_1) < T$). Let y be the analogous vertex in P_2 (and let $y = t$ if $\ell(P_2) < T$). We say that P_v *passes the T-test* if the portion of P_v between x and y is a shortest path. See Figure 3 for an example.

LEMMA 4.1. *If P_v passes the T-test, then P_v is T-LO.*

PROOF. Suppose P_v passes the test and consider a subpath P' of P_v as in the definition of T-LO. If P' is a subpath of P_1 or P_2, then it is a shortest path. Otherwise P' contains v and is a subpath of the portion of P_v between x and y (as defined in the T-test), and therefore also a shortest path. □

This test is very efficient: it traverses P_v at most once and runs a single point-to-point shortest-path query. Although it may miss some admissible paths (a T-LO path may fail the T-test), it can be off by at most a factor of two:

LEMMA 4.2. *If P_v fails the T-test, then P_v is not $2T$-LO.*

PROOF. If P_v fails the test, then the x–y subpath in the definition of the test is not a shortest path. Delete x and y from the subpath, creating a new path P''. We know

$\ell(P'') < 2T$ (v divides it into two subpaths of length less than T), which means P_v is not $2T$-LO. □

Computing the exact UBS of a single via path P_v in a straightforward manner requires $O(|P_v|^2)$ point-to-point queries, which is too slow for practical applications. The following lemma allows us to bound the UBS of P_v by considering only the stretch of the path, which is trivial to compute.

LEMMA 4.3. *If a via path P_v has stretch $(1+\epsilon)$ and passes the T-test for $T = \beta \cdot \text{dist}(s,t)$ (with $0 < \epsilon < \beta < 1$), then P_v is a $\frac{\beta}{\beta-\epsilon}$-UBS path.*

PROOF. Consider a subpath P' of P_v between vertices u and w. If $v \notin P'$, or both u and w are within distance T of v, then P' is a shortest path (as a subpath of a shortest path). Assume v is between u and w and at least one of these vertices is at distance more than T from v. This implies $\ell(P') \geq T = \beta \cdot \text{dist}(s,t)$. Furthermore, we know that $\ell(P') \leq \text{dist}(u,w) + \epsilon \cdot \text{dist}(s,t)$ (the absolute stretch of the subpath P' cannot be higher than in P_v). Combining these two inequalities, we get that $\ell(P') \leq \text{dist}(u,w) + \epsilon \cdot \ell(P')/\beta$. Rearranging the terms, we get that $\ell(P') \leq \beta \cdot \text{dist}(u,w)/(\beta - \epsilon)$, which completes the proof. □

4.1. BDV Algorithm

We now consider a relatively fast BD-based algorithm, which we call BDV. It grows shortest path trees from s and into t; each search stops when it advances more that $(1+\epsilon)\ell(Opt)$ from its origin. (This is the longest an admissible path can be.) For each vertex v scanned in both directions, we check whether the corresponding path P_v is *approximately admissible*: it shares at most $\gamma \cdot \ell(Opt)$ with Opt, has limited stretch ($\ell(P_v) \leq (1+\epsilon)\ell(Opt)$), and passes the T-test for $T = \alpha \cdot \ell(Opt)$. Finally, we output the best approximately admissible via path according to the objective function.

4.2. The Choice Routing Algorithm

A related method is the *choice routing algorithm* (CR) [Cambridge Vehicle Information Technology Ltd. 2005]. It starts by building shortest path trees from s and to t. It then looks at *plateaus*, i.e., maximal paths that appear in both trees simultaneously. In general, a plateau u–w gives a natural s–t path: follow the out-tree from s to u, then the plateau, then the in-tree from w to t. The algorithm selects paths corresponding to long plateaus, orders them according to some "goodness" criteria (not explained in the original paper [Cambridge Vehicle Information Technology Ltd. 2005]), and outputs the best one (or more, if desired). Because the paths found by CR have large plateaus, they have good local optimality:

LEMMA 4.4. *If P corresponds to a plateau v–w, P is $\text{dist}(v,w)$-LO.*

PROOF. If P is not a shortest path, then there are vertices x, y on P such that the length of P between these vertices exceeds $\text{dist}(x,y)$. Then x must strictly precede v on P, and y must strictly follow w. This implies the lemma. □

Note that both CR and BDV are based on BD and only examine single-via paths. While BDV must run one point-to-point query to evaluate each candidate path, all plateaus can be detected in linear time. This means CR has the same complexity as BD (ignoring the time for goodness evaluation), which is much faster than BDV. It should be noted, however, that local optimality can be achieved even in the absence of long plateaus (as illustrated in Appendix A). This means BDV can potentially succeed in situations where CR fails. Still, neither method is fast enough for continental-sized road networks.

5. PRUNING

A natural approach to accelerate BD is to prune it at unimportant vertices (as done by RE or CH, for example). In this section, we show how known pruning algorithms can be extended to find admissible single-via paths. The results of Abraham, Fiat et al. [2010b] suggest that pruning is unlikely to discard promising via vertices. Because of local optimality, an admissible alternative path P contains a long shortest subpath P' that shares little with Opt. Being a shortest path, P' must contain an "important" (unpruned) vertex v.

For concreteness, we focus on an algorithm based on RE; we call it REV. Like BDV, REV builds two (now pruned) shortest paths trees, out of s and into t. We then evaluate each vertex v scanned by both searches as follows. First, we perform two P2P queries (s–v and v–t) to find P_v. (They are necessary because some original tree paths may be suboptimal due to pruning.) We then perform an approximate admissibility test on P_v, as in BDV. Among all candidate paths that pass, we return the one minimizing the objective function $f(\cdot)$.

The main advantage of replacing BD by RE is a significant reduction in the number of via vertices we consider. Moreover, auxiliary P2P queries (including T-tests) can also use RE, making REV potentially much faster than BDV.

An issue we must still deal with is computing the sharing amount $\sigma(v)$. RE adds shortcuts to the graph, each representing a path in the original graph. To calculate $\sigma(v)$ correctly, we must solve a *partial unpacking* subproblem: given a shortcut (a, c), with $a \in Opt$ and $c \notin Opt$, find the vertex $b \in Opt$ that belongs to the a–c path and is farthest from a. Assuming each shortcut bypasses exactly one vertex and the shortcut hierarchy is balanced (as is usually the case in practice), this can be done in $O(\log n)$ time with binary search.

5.1. Running Time

We can use the results of Abraham, Fiat et al. [2010b] to analyze REV. As required in their paper, the analysis assumes the input network is undirected, but our algorithms and implementations work just as well on directed graphs.

Suppose we have a constant-degree network of diameter D and highway dimension h. The reach-based query algorithm scans $O(k \log D)$ vertices and runs in time $O((k \log D)^2)$, where k is either h or $h \log n$, depending on whether preprocessing must be polynomial or not. For each vertex v scanned in both directions, REV needs a constant number of P2P queries and partial unpackings. The total running time is therefore $O((k \log D)^3 + k \log D \log n)$, which is sublinear (unlike BDV or CR). The same algorithm (and analysis) can be applied if we use contraction hierarchies; we call the resulting algorithm CHV.

5.2. Relaxed Reaches

Pruning may cause REV and CHV to miss candidate via vertices, leading to suboptimal solutions. In some cases, they may not find any admissible path even when BDV would. For REV, we can fix this by trading off some efficiency. If we multiply the reach values by an appropriate constant, the algorithm is guaranteed to find all admissible single-via paths. The resulting algorithm is as effective as BDV, but much more efficient. It exploits the fact that vertices in the middle of locally optimal paths have high reach:

LEMMA 5.1. *If the path P is T-LO and $v \in P$, then* $r(v) \geq \min\{T/2, \text{dist}(s, v), \text{dist}(v, t)\}$.

PROOF. Let v be at least $T/2$ away from the endpoints of P. Let x and y be the closest vertices to v that are at least $T/2$ away from v towards s and t, respectively.

Since P is T-LO, the subpath of P between x and y is a shortest path, and v has reach at least $T/2$. □

COROLLARY 5.2. *If the path P passes the T-test and $v \in P$, then $r(v) \geq \min\{T/2, \text{dist}(s,v), \text{dist}(v,t)\}$.*

Let δ-REV be a version of REV that uses original reach values multiplied by $\delta \geq 1$ to prune the original trees from s and to t (i.e., it uses $\delta \cdot r(v)$ instead of $r(v)$). Auxiliary P2P computations to build and test via paths can still use the original $r(v)$ values. The algorithm clearly remains correct, but may prune fewer vertices. The parameter δ gives a trade-off between efficiency and success rate:

THEOREM 5.3. *If $\delta \geq 2(1+\epsilon)/\alpha$, δ-REV finds the same admissible via paths as* BDV.

PROOF. Consider a via path P_v that passes the T-test for $T = \alpha \cdot \text{dist}(s,t)$. Then by Corollary 5.2 for every vertex $u \in P_v$, $r(u) \geq \min\{\text{dist}(s,v), \text{dist}(v,t), \alpha \cdot \text{dist}(s,t)/2\}$. When $\delta \geq 2(1+\epsilon)/\alpha$, then $\delta \cdot r(u) \geq \min\{\text{dist}(s,v), \text{dist}(v,t), (1+\epsilon)\text{dist}(s,t)\}$. Therefore no vertex on P_v is pruned, implying that $\text{dist}(s,v)$ and $\text{dist}(v,t)$ are computed correctly. As a result P_v is considered as an admissible via path. □

The analysis of Abraham, Fiat et al. [2010b] implies that multiplying reach values by a constant increases the query complexity by a constant multiplicative factor. Hence, the asymptotic time bounds for REV also apply to δ-REV, for any constant $\delta \geq 1$. Note that Theorem 5.3 assumes that δ-REV and BDV are applied on the same graph, with no shortcuts.

5.3. Relaxed Contraction Hierarchies

Because CH is more efficient than RE, CHV prunes the search space even more than REV does. Although this leads to better query times, the reduced number of candidate via paths may lead to lower solution quality. Similar to the relaxation of REV, we can also relax CHV by allowing it to "look downwards" (to lower-order vertices) during the query.

The details are as follows. During a query, let $p_i(u)$ be the ith ancestor of u in the search tree: $p_1(u)$ is u's parent, and $p_i(u)$ is the parent of $p_{i-1}(u)$ (for $i > 1$). The k-relaxed variant of CHV prunes an edge (u,v) only if v precedes all vertices u, $p_1(u), \ldots, p_k(u)$ in the CH order. If u has fewer than k ancestors, (u,v) is never pruned.

Note that the relaxed variant of REV is somewhat more natural and has a theoretical justification. As we shall see, however, the proposed relaxation of CHV works, allowing a trade-off between running times and solution quality.

6. PRACTICAL ALGORITHMS

The algorithms proposed so far produce a set of candidate via paths and explicitly check whether each is admissible. We introduced techniques to reduce the number of candidates and to check (approximate) admissibility faster, with a few point-to-point queries. Unfortunately, this is not enough in practice. Truly practical algorithms can afford at most a (very small) constant number of point-to-point queries *in total* to find an alternative path. Therefore, instead of actually evaluating all candidate paths, in our experiments we settle for finding one that is good enough. As we will see, we sort the candidate paths according to some objective function, test the vertices in this order, and return the first admissible path as our answer. We consider two versions of this algorithm, one using BD and the other (the truly practical one) a pruned shortest path algorithm (RE or CH). Although based on the methods we introduced in previous sections, the solutions they find do not have the same theoretical guarantees. In par-

ticular, they may not return the best (or any) via path, even if one exists. As Section 7 will show, however, the heuristics still have very high success rates in practice.

The practical algorithms sort the candidate paths P_v in nondecreasing order according to the function $f(v) = 2\ell(v)+\sigma(v)-pl(v)$, where $\ell(v)$ is the length of the via path P_v, $\sigma(v)$ is how much P_v shares with Opt, and $pl(v)$ is the length of the longest plateau containing v. Note that $pl(v)$ is a lower bound on the local optimality of P_v (by Lemma 4.4); by preferring vertices with high $pl(v)$, we tend to find an admissible path sooner.

6.1. X-BDV

We are now ready to describe X-BDV, an experimental version of BDV that incorporates elements of CR for efficiency. Although much slower than our pruning algorithms, this version is fast enough to run experiments on (unlike BDV). It runs BD, with each search stopping when its radius is greater than $(1+\epsilon)\ell(Opt)$; we also prune any vertex u with $\text{dist}(s,u)+\text{dist}(u,t) > (1+\epsilon)\ell(Opt)$. In linear time, we compute $\ell(v)$, $\sigma(v)$, and $pl(v)$ for all vertices v visited by both searches. We return the path P_v with highest $f(v)$ value among all paths that satisfy three hard constraints: $\ell(P_v \setminus Opt) < (1+\epsilon)\ell(Opt \setminus P_v)$ (the detour is not much longer than the subpath it skips), $\sigma(v) < \gamma \cdot \ell(Opt)$ (sharing is limited), and $pl(v) > \alpha \cdot \ell(P_v \setminus Opt)$ (the plateau is large, implying enough local optimality). Note that we specify local optimality relative to the detour only (and not the entire path, as in Section 3). Our rationale for doing so is as follows. In practice, alternatives sharing up to 80% with *Opt* may still make sense. In such cases, the T-test will always fail unless the (path-based) local optimality is set to 10% or less. This is too low for alternatives that share nothing with *Opt*. Using detour-based local optimality is a reasonable compromise. For consistency, ϵ is also used to bound the stretch of the detour (as opposed to the entire path).

6.2. Pruned Algorithms: X-REV and X-CHV

The second algorithm we tested, X-REV, is similar to X-BDV but grows reach-pruned trees out of s and into t. The stopping criteria and the evaluation of each via vertex v are the same as in X-BDV. As explained in Section 4, RE trees may give only upper bounds on $\text{dist}(s,v)$ and $\text{dist}(v,t)$ for any vertex $v \notin Opt$. But we can still use the approximate values (given by the RE trees) of $\ell(v)$, $\sigma(v)$, and $pl(v)$ to sort the candidate paths in nondecreasing order according to $f(v)$. We then evaluate each vertex v in this order as follows. We first compute the actual via path P_v with two RE queries, s–v and v–t (as an optimization, we reuse the original forward tree from s and the backward tree from t). Then we compute the exact values of $\ell(v)$ and $\sigma(v)$ and check whether $\ell(P_v \setminus Opt) < (1+\epsilon)\ell(Opt \setminus P_v)$, whether $\sigma(v) < \gamma \cdot \ell(Opt)$, and run a T-test with $T = \alpha \cdot \ell(P_v \setminus Opt)$. If P_v passes all three tests, we pick it as our alternative. Otherwise, we discard v as a candidate, penalize its descendants (in both search trees) and try the next vertex in the list. We penalize a descendant u of v in the forward (backward) search tree by increasing $f(u)$ by $\text{dist}(s,v)$ ($\text{dist}(v,t)$). This gives less priority to vertices that are likely to fail, keeping the number of check queries small.

Of course, in order to improve the success rate of this algorithm, we can apply the relaxed reach value approach from Section 5 to X-REV as well.

A third implementation we tested was X-CHV, which is similar to X-REV but uses contraction hierarchies (rather than reaches) for pruning. We study both the pure and the relaxed variants of this algorithm.

7. EXPERIMENTS

We implemented the algorithms from Section 6 in C++ and compiled them with Microsoft Visual C++ 2008. Queries use a binary heap as priority queue. The evaluation was conducted on a dual AMD Opteron 250 running Windows 2003 Server. It is

clocked at 2.4 GHz and has 16 GB of RAM and 2 x 1 MB of L2 cache. Our code is single-threaded and runs on a single processor at a time.

7.1. Methodology

We use the European road network, with 18 million vertices and 42 million edges, made available for the 9th DIMACS Implementation Challenge [Demetrescu et al. 2006]. It uses travel times as the length function. (We also experimented with TIGER/USA data, but closer examination revealed that the data has errors, with missing arcs on several major highways and bridges. This makes the results less meaningful, so we do not include them.) We allow the detour to have maximum stretch $\epsilon = 25\%$, set the maximum sharing value to $\gamma = 80\%$, and set the minimum detour-based local optimality to $\alpha = 25\%$ (see Section 6).

X-REV and X-CHV extend the point-to-point query algorithms RE [Goldberg et al. 2009] and CH [Geisberger et al. 2008]. Since preprocessing does not change, we use the preprocessed data given in the original papers [Geisberger et al. 2008; Goldberg et al. 2009]. In particular, preprocessing takes 45 minutes for RE and 25 for CH. Note that the reach values we use (computed by Goldberg et al. [2009]) are upper bounds; computing exact reaches is prohibitively expensive.

We compare the algorithms in terms of both query performance and path quality. Performance is measured by the number of vertices scanned and by query times (given in absolute terms and as a multiple of the corresponding P2P method). Quality is given first by the *success rate*: how often the algorithm finds as many alternatives as desired. Among the successful runs, we also compute the average and worst uniformly bounded stretch, sharing, and detour-based local optimality. (Reporting these values requires $O(|P|^2)$ point-to-point queries for each path P; this evaluation is not included in the query times.) Unless otherwise stated, figures are based on 1 000 queries, with source s and target t chosen uniformly at random.

7.2. Success Rate

In our first experiment, reported in Table I, we vary p, the desired number of alternatives to be computed by the algorithms. There is a clear trade-off between success rates and query times. As expected, X-BDV is successful more often, while X-CHV is fastest. Unfortunately, X-BDV takes more than half a minute for each query and X-CHV finds an alternative in only 58.2% of the cases (the numbers are even worse with two or three alternatives). The reach-based algorithm, X-REV, seems to be a good compromise between these extremes. Queries are still fast enough to be practical, and it is almost as successful as X-BDV. In only 20.4 ms, it finds a good alternative in 91.3% of the cases.

Alternative paths, when found, tend to have similar quality, regardless of the algorithm. This may be because the number of admissible alternatives is small: the algorithms are much more successful at finding a single alternative than at finding three (see Table I). On average, the first alternative has 10% stretch, is 72% locally optimal, and shares around 47% with the optimum. Depending on the success rate and p, the alternative query algorithm is 4 to 12 times slower than a simple P2P query with the same algorithm. This is acceptable, considering how much work is required to identify good alternatives.

7.3. Relaxed Pruning

As observed in Section 6, we can increase the success rates of X-REV and X-CHV by relaxing the pruning rules during the query. For X-REV we multiply the reach values by a constant $\delta > 1$, whereas k-relaxed X-CHV looks "downwards" during the query. Our second experiment evaluates how these relaxations affect both algorithms. Table II re-

Table I. Performance of various algorithms on the European road network as the number of desired alternatives (p) changes. Column *success rate* reports how often the algorithm achieves this goal. For the successful cases, we report the (average and worst-case) quality of the p-th alternative in terms of UBS, sharing, and detour-based local optimality. Finally, we report the average number of scanned vertices and query times (both in milliseconds and as a multiple of the corresponding point-to-point variant).

		PATH QUALITY							PERFORMANCE		
		success	UBS[%]		sharing[%]		loc opt[%]		#scanned	time	slow-
p	algo	rate[%]	avg	max	avg	max	avg	min	vertices	[ms]	down
1	X-BDV	94.5	9.4	35.8	47.2	79.9	73.1	30.3	16 963 507	26 352.0	6.0
	X-REV	91.3	9.9	41.8	46.9	79.9	71.8	30.7	16 111	20.4	5.6
	X-CHV	58.2	9.6	45.8	42.9	79.9	74.3	30.6	1 510	3.1	4.6
2	X-BDV	81.1	11.8	38.5	62.4	80.0	71.8	29.6	16 963 507	29 795.0	6.8
	X-REV	70.3	12.2	38.1	60.3	80.0	71.3	29.6	25 322	33.6	9.2
	X-CHV	28.6	10.8	45.4	55.3	79.6	77.6	30.3	1 685	3.6	5.3
3	X-BDV	61.6	13.2	41.2	68.9	80.0	68.7	30.6	16 963 507	33 443.0	7.7
	X-REV	43.0	12.8	41.2	66.6	80.0	74.9	33.3	30 736	42.6	11.7
	X-CHV	10.9	12.0	41.4	59.3	80.0	79.0	36.1	1 748	3.9	5.8

ports the success rate and query times of δ-X-REV (k-relaxed X-CHV) for several values of δ (k) when $p = 1$, while Table III gives the corresponding results for $p = 2$ and $p = 3$.

As predicted, for $p = 1$, higher reach bounds do improve the success rate, eventually matching that of X-BDV on average. The worst-case UBS is also reduced from almost 42% to around 30% when δ increases. Furthermore, most of the quality gains are already obtained with $\delta = 2$, when queries are still fast enough (less than 10 times slower than a comparable P2P query). For X-CHV, relaxing the pruning rule clearly helps: setting $k = 3$ increases the poor success rate of the non-relaxed version to satisfactory levels.

For $p = 2$ and $p = 3$, the results are similar to those of finding the first path but more pronounced, especially for the third path. Increasing δ from 1 to 2 yields the biggest improvement in the success rate of X-REV. The improvement is still noticeable for $\delta = 3$, but reaches the point of diminishing returns thereafter. For $\delta = 4$, the success rate is similar to that of X-BDV, while for $\delta = 10$, X-REV finds an alternative more often than X-BDV. (Recall that, unlike X-REV, X-BDV discards alternative paths based on plateau lengths only, and therefore may miss admissible alternatives considered by

Table II. Performance of X-REV and X-CHV when relaxing the pruning rules. Column δ, k indicates the multiplier for reach values and how much we relax X-CHV. We set $p = 1$.

		PATH QUALITY							PERFORMANCE		
		success	UBS[%]		sharing[%]		loc opt[%]		#scanned	time	slow-
algo	δ, k	rate[%]	avg	max	avg	max	avg	min	vertices	[ms]	down
X-REV	1	91.3	9.9	41.8	46.9	79.9	71.8	30.7	16 111	20.4	5.6
	2	94.2	9.7	31.6	46.6	79.9	71.3	27.6	31 263	34.3	9.4
	3	94.2	9.5	29.2	46.7	79.9	71.9	31.2	53 464	55.3	15.2
	4	94.3	9.5	29.3	46.7	79.9	71.8	31.2	80 593	83.2	22.8
	5	94.4	9.5	29.3	46.7	79.9	71.8	31.4	111 444	116.6	31.9
	10	94.6	9.5	30.2	46.8	79.9	71.7	31.4	289 965	344.3	94.3
X-CHV	0	58.2	9.6	45.8	42.9	79.9	74.3	30.6	1 510	3.1	4.6
	1	80.0	10.8	63.4	46.9	80.0	70.6	27.6	3 652	6.2	9.1
	2	87.5	11.5	52.5	45.8	80.0	67.8	26.2	6 756	9.4	13.9
	3	90.7	11.5	45.4	45.4	80.0	67.7	30.0	12 104	16.9	25.0
	4	92.0	11.3	41.9	45.2	80.0	67.1	27.9	22 025	29.2	43.3
	5	92.9	11.3	41.9	44.8	80.0	66.8	27.9	39 835	55.2	81.8
	6	93.7	11.1	41.9	44.3	80.0	66.9	27.9	71 098	105.2	155.9
	8	94.3	11.0	41.9	44.0	80.0	66.8	27.9	210 046	368.8	546.4
	10	94.7	11.0	47.7	43.6	80.0	66.4	26.3	558 516	1 225.6	1 815.7
X-BDV	–	94.5	9.4	35.8	47.2	79.9	73.1	30.3	16 963 507	26 352.0	6.0

Table III. Performance of X-REV and X-CHV for $p = 2$ and $p = 3$ when relaxing the pruning rules.

		$p = 2$				$p = 3$			
algo	δ, k	success rate[%]	#scanned vertices	time [ms]	slow-down	success rate[%]	#scanned vertices	time [ms]	slow-down
X-REV	1	70.3	25 322	33.6	9.2	43.0	30 736	42.6	11.7
	2	79.0	41 259	50.3	13.8	56.9	48 688	64.9	17.8
	3	80.6	63 526	73.0	20.0	61.0	71 399	91.7	25.2
	4	81.7	90 243	103.3	28.3	61.6	98 661	126.4	34.7
	5	81.7	121 073	142.9	39.1	61.6	129 890	166.2	45.6
	10	81.9	299 042	392.9	107.6	62.6	310 004	446.2	122.5
X-CHV	0	28.6	1 685	3.6	5.3	10.9	1 748	3.7	5.5
	1	49.5	3 985	7.5	11.1	24.0	4 125	7.6	11.3
	2	61.6	7 230	11.9	17.6	34.7	7 463	13.1	19.4
	3	70.1	12 690	20.3	30.1	42.3	13 011	22.1	32.7
	4	74.0	22 708	35.3	52.3	48.5	23 135	37.8	56.0
	5	77.0	40 657	65.0	96.3	53.3	41 185	73.2	108.4
	6	78.9	72 047	119.8	177.5	55.8	72 680	135.6	200.9
	8	80.7	211 323	433.8	642.7	59.5	212 247	484.1	717.2
	10	81.9	560 221	1 445.1	2 140.9	61.8	561 563	1 614.4	2 391.7
X-BDV	–	81.1	16 963 507	29 795.0	6.8	61.6	16 963 507	33 443.0	7.7

X-REV.) For X-CHV, increasing k up to 10 helps the success rate, but the loss in query performance does not justify increases beyond $k = 6$.

Comparing X-CHV and X-REV, we observe that X-CHV relaxed with $k = 3$ achieves the same success rate as non-relaxed X-REV, but queries are 1.2 ($p = 1$) to 1.9 ($p = 3$) times faster. However, increasing k further is not as effective as increasing δ. Compared to relaxed X-REV with $\delta = 2$, we have to increase k to 6 to achieve similar success rates with X-CHV. For these parameter values, X-CHV is more than twice as slow. We conclude that when the emphasis is on query speed, 3-relaxed X-CHV is the best choice, whereas 2-relaxed X-REV finds better solutions at the expense of higher running times.

The original reach values used by X-REV are not exact: they are upper bounds computed by a particular preprocessing algorithm [Goldberg et al. 2009]. On a smaller graph (of the Netherlands, with $n \approx 0.9$M and $m \approx 2.2$M), we could actually afford to compute exact reaches on the shortcut-enriched graph output by the standard RE preprocessing. On this graph, the success rate drops from 83.4% with the original upper bounds to 81.7% with exact reaches. Multiplying the exact reaches by $\delta = 2$ increases the success rate again to 83.6%. With $\delta = 5$, we get 84.7%, very close to the 84.9% obtained by X-BDV. Note that the Dutch subgraph tends to have fewer alternatives than Europe as a whole.

7.4. Local Queries

Up to now, we have only considered random queries, which are mostly long-distance. In real-world applications, however, many queries are local. In order to evaluate such queries, Figure 4 reports the success rate of our algorithms for $p = 1, 2, 3$ and various Dijkstra ranks on Europe. The *Dijkstra rank* of v with respect to s is i if v is the ith vertex taken from the priority queue when running a Dijkstra query from s. The results are based on $1\,000$ queries for each rank. We observe that the success rate is lower for local queries, as expected. Still, for mid-range queries we find an alternative in 60% to 80% of the cases, which seems reasonable. (Recall that an admissible alternative may not exist.) Multiplying reach values helps all types of queries: 2-X-REV has almost the same success rate as X-BDV for all ranks and number of alternatives examined. These observations also hold for the second and third alternatives. We observe that X-CHV finds fewer alternatives than any other algorithm, independent of the Dijkstra rank or the number of alternatives, but this can be (partly) remedied by relaxing it with $k = 3$.

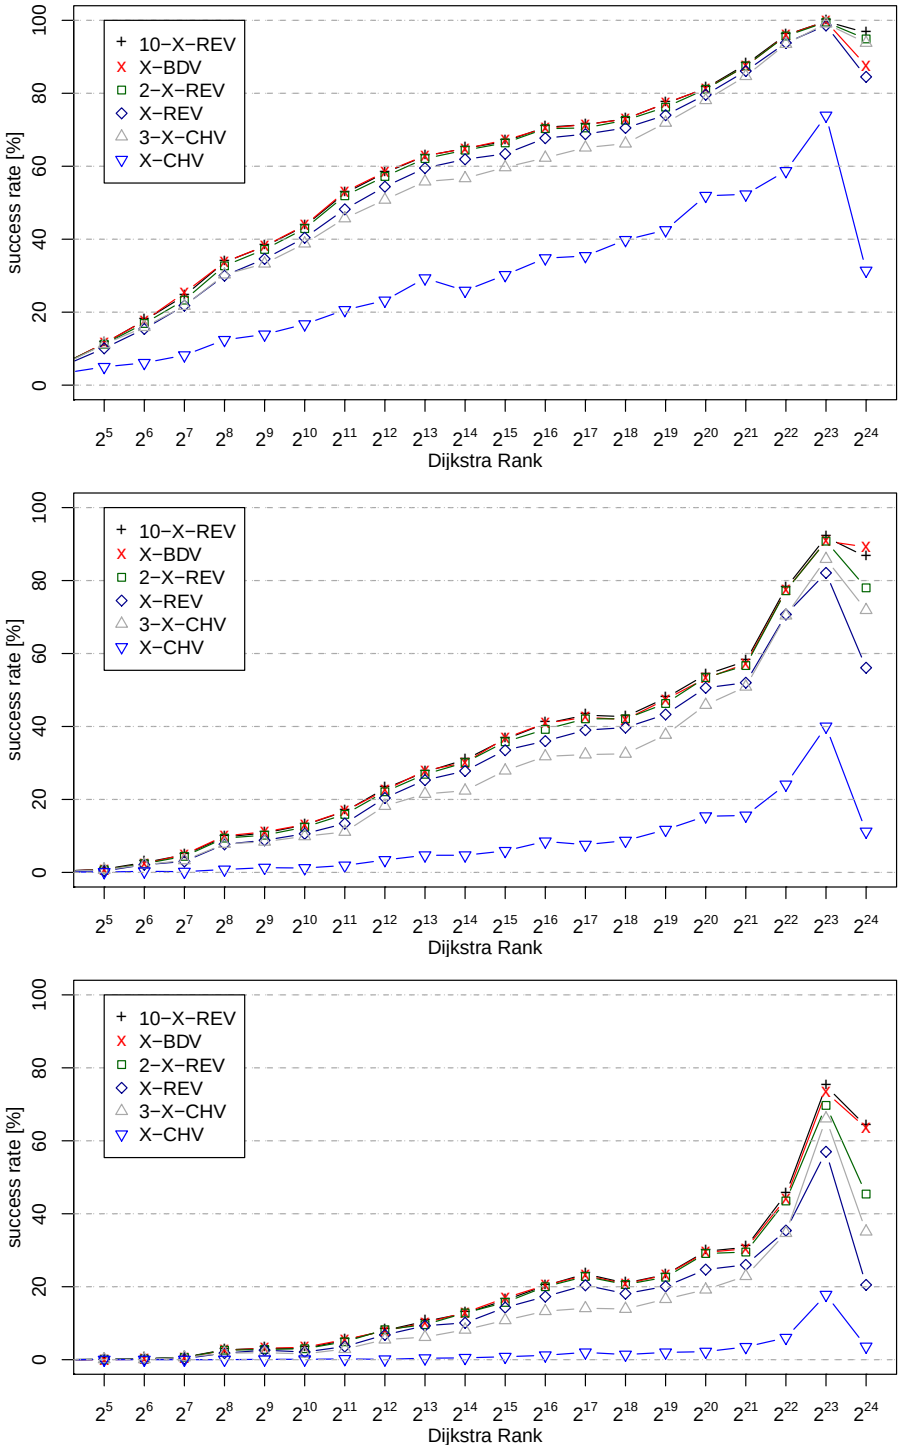

Fig. 4. Success rate for X-BDV, X-REV, 2-X-REV, 10-X-REV, X-CHV, and 3-relaxed X-CHV when varying the Dijkstra rank of queries for $p = 1$ (top), $p = 2$ (middle), and $p = 3$ (bottom).

Table IV. Performance of the iterative approach compared to X-BDV. The columns refer to the same values as explained in Tab. I.

	PATH QUALITY							PERFORMANCE		
	success	UBS[%]		sharing[%]		loc opt[%]		#scanned	time	slow-
algo	rate[%]	avg	max	avg	max	avg	min	vertices	[ms]	down
X-BDV	94.5	9.4	35.8	47.2	79.9	73.1	30.3	16 963 507	26 352.0	6.0
BDI	95.1	4.9	21.6	28.5	78.7	4.4	0.0	10 042 588	8 952.3	2.0
BDI-2	51.7	76.4	551.2	48.4	78.1	7.5	0.0	25 467 501	22 573.0	5.2

7.5. Iterative Approach

Our experiments so far have omitted a natural approach to computing alternative paths, which we discuss now for completeness. It works by first computing the shortest path *Opt*. Then it sets $P = Opt$ and it the repeats the following steps:

(1) multiply the current length of every edge in P by β;
(2) recompute the shortest path Opt' from s to t in the graph G' with the updated edge lengths;
(3) if Opt' fulfills certain constraints, return Opt' as an alternative; otherwise set $P \leftarrow Opt'$ and go to step 1.

Note that we have to determine $\ell(Opt')$ with respect to G since it was computed on a graph with modified edge lengths. We implemented a version of this algorithm based on bidirectional Dijkstra; we call it BDI. We use $\beta = 1.2$. As in previous experiments, we accept a path if it shares less than 80% with *Opt* and if it is not more than 25% longer than *Opt*.

Incorporating a constraint for "reasonable" paths is harder, however, since there is no natural way of running a small number of T-tests to ensure a minimum degree of local optimality (violations of local optimality could appear anywhere on the path). This is in fact the main drawback of this approach: all we can do is run the algorithm and hope the paths it produces will be acceptable.

Another issue with the iterative-based approach is that it often finds alternatives with large numbers of local detours, which intuitively is not what users expect. To avoid particularly bad cases, we also tested BDI-2, a variant of BDI with one additional constraint: only accept alternatives with at most two detours from the shortest path. Note that this is still not as restrictive as the via-vertex algorithms, which are allowed a single detour.

Table IV compares BDI and BDI-2 to X-BDV. The algorithms are evaluated according to the same criteria as before: success rate, uniformly bounded stretch, sharing, local optimality, and running times.

At first sight, the results are promising. Compared to X-BDV BDI has higher success rate (95.1%), smaller stretch, and lower sharing. Moreover, X-BDV is faster: in most cases, it finds a valid alternative after penalizing the optimal path only once. Unfortunately, the alternatives provided by BDI have a fundamental drawback: the local optimality is extremely low. A related issue (not shown in the table) is that alternatives have more than 5 detours on average, and up to 15 in the worst case. Such alternatives would seem unnatural to most users. Restricting the number of detours to at most two (BDI-2) does not help much: it improves local optimality only slightly, and success rates and sharing are significantly worse. We conclude that X-BDV is superior to BDI and BDI-2.

Another shortcoming of the iterative approach is its combination with RE or CH. After increasing the lengths of a path, we have to update the preprocessed data. While this is doable (but time-consuming) for CH [Schultes and Sanders 2007], it is unknown how to update reach values efficiently. Of course, one can use heuristics (such as pe-

nalizing only the shortcuts on the path, not the original edges) but preliminary experiments showed that the resulting path quality is very bad.

We stress that there may be cleverer ways of implementing the penalty-based approach [Bader et al. 2011; Chen et al. 2007]. Still, the via-vertex approach is only about 6 times slower than its point-to-point variant, which means the iterative approach would have to find an alternative after at most 6 iterations to be competitive. Moreover, guaranteeing good local optimality with the iterative approach seems hard.

8. CONCLUSION

By introducing the notion of admissibility, we have given the first formal treatment to the problem of finding alternative paths. The natural concept of local optimality allows us to prove properties of such paths. Moreover, we have given theoretically efficient algorithms for an important subclass, that of single via paths. We concentrated on making the approach efficient for real-time applications by designing approximate admissibility tests and an optimization function biased towards admissible paths. Our experiments have shown that these simplified versions are practical for real, continental-sized road networks.

More generally, however, our techniques allow us to do optimization constrained to the (polynomially computable) set of admissible single via paths. We could optimize other functions over this set, such as fuel consumption, time in traffic, or tolls. This gives a heuristic alternative to multi-criteria optimization.

Our work leads to natural open questions. In particular, are there efficient exact tests for local optimality and uniformly bounded stretch? Furthermore, can one find admissible paths with multiple via vertices efficiently? This is especially interesting because it helps computing arbitrary admissible paths, since any admissible alternative with stretch $1+\epsilon$ and local optimality $\alpha \cdot \ell(Opt)$ is defined by at most $\lceil (1+\epsilon)/\alpha \rceil - 1$ via points.

Although our code is purely sequential, there are several natural opportunities for parallelization. Besides using two cores for forward and backward search, we can parallelize the evaluation of candidates (reconstructing the real path and running the T-test). We could use multiple cores to evaluate several alternatives at once and pick the best, which would lead to better alternatives in essentially the same time.

REFERENCES

ABRAHAM, I., DELLING, D., GOLDBERG, A. V., AND WERNECK, R. F. 2010a. Alternative Routes in Road Networks. In *Proceedings of the 9th International Symposium on Experimental Algorithms (SEA'10)*, P. Festa, Ed. Lecture Notes in Computer Science Series, vol. 6049. Springer, 23–34.

ABRAHAM, I., FIAT, A., GOLDBERG, A. V., AND WERNECK, R. F. 2010b. Highway Dimension, Shortest Paths, and Provably Efficient Algorithms. In *Proceedings of the 21st Annual ACM–SIAM Symposium on Discrete Algorithms (SODA'10)*. 782–793.

BADER, R., DEES, J., GEISBERGER, R., AND SANDERS, P. 2011. Alternative Route Graphs in Road Networks. In *Proceedings of the 1st International ICST Conference on Theory and Practice of Algorithms in (Computer) Systems (TAPAS'11)*, A. Marchetti-Spaccamela and M. Segal, Eds. Lecture Notes in Computer Science Series, vol. 6595. Springer, 21–32.

CAMBRIDGE VEHICLE INFORMATION TECHNOLOGY LTD. 2005. Choice Routing. Available at `http://www.camvit.com`.

CHEN, Y., BELL, M. G. H., AND BOGENBERGER, K. 2007. Reliable Pretrip Multipath Planning and Dynamic Adaptation for a Centralized Road Navigation System. *IEEE Transactions on Intelligent Transportation Systems 8,* 1, 14–20.

DANTZIG, G. B. 1962. *Linear Programming and Extensions*. Princeton Univ. Press, Princeton, NJ.

DELLING, D., SANDERS, P., SCHULTES, D., AND WAGNER, D. 2009. Engineering Route Planning Algorithms. In *Algorithmics of Large and Complex Networks*, J. Lerner, D. Wagner, and K. Zweig, Eds. Lecture Notes in Computer Science Series, vol. 5515. Springer, 117–139.

DELLING, D. AND WAGNER, D. 2009. Pareto Paths with SHARC. In *Proceedings of the 8th International Symposium on Experimental Algorithms (SEA'09)*, J. Vahrenhold, Ed. Lecture Notes in Computer Science Series, vol. 5526. Springer, 125–136.

DEMETRESCU, C., GOLDBERG, A. V., AND JOHNSON, D. S., Eds. 2006. *9th DIMACS Implementation Challenge - Shortest Paths*.

DIJKSTRA, E. W. 1959. A Note on Two Problems in Connexion with Graphs. *Numerische Mathematik 1*, 269–271.

EPPSTEIN, D. 1994. Finding the k shortest paths. In *Proceedings of the 35th Annual IEEE Symposium on Foundations of Computer Science (FOCS'94)*. 154–165.

GEISBERGER, R., KOBITZSCH, M., AND SANDERS, P. 2010. Route Planning with Flexible Objective Functions. In *Proceedings of the 12th Workshop on Algorithm Engineering and Experiments (ALENEX'10)*. SIAM, 124–137.

GEISBERGER, R., SANDERS, P., SCHULTES, D., AND DELLING, D. 2008. Contraction Hierarchies: Faster and Simpler Hierarchical Routing in Road Networks. In *WEA'08*, C. C. McGeoch, Ed. Lecture Notes in Computer Science Series, vol. 5038. Springer, 319–333.

GOLDBERG, A. V., KAPLAN, H., AND WERNECK, R. F. 2009. Reach for A*: Shortest Path Algorithms with Preprocessing. In *The Shortest Path Problem: Ninth DIMACS Implementation Challenge*, C. Demetrescu, A. V. Goldberg, and D. S. Johnson, Eds. DIMACS Book Series, vol. 74. American Mathematical Society, 93–139.

GUTMAN, R. J. 2004. Reach-Based Routing: A New Approach to Shortest Path Algorithms Optimized for Road Networks. In *Proceedings of the 6th Workshop on Algorithm Engineering and Experiments (ALENEX'04)*. SIAM, 100–111.

HANSEN, P. 1979. Bricriteria Path Problems. In *Multiple Criteria Decision Making: Theory and Application*, G. Fandel and T. Gal, Eds. Springer, 109–127.

MARTINS, E. Q. 1984. On a Multicriteria Shortest Path Problem. *European Journal of Operational Research 26,* 3, 236–245.

SCHULTES, D. AND SANDERS, P. 2007. Dynamic Highway-Node Routing. In *Proceedings of the 6th Workshop on Experimental Algorithms (WEA'07)*, C. Demetrescu, Ed. Lecture Notes in Computer Science Series, vol. 4525. Springer, 66–79.

Received November 2010; revised March 2011; accepted September 2012

A. APPENDIX

This appendix contains additional examples of the output of X-REV with Europe as input, shown in Figure 6).

In addition, we illustrate how BDV may find alternatives that are pruned by CR. Consider the example given in Figure 5. The shortest path is $Opt = (s, a, t)$. For appropriate choice of parameters, admissible alternative paths are $B = (s, e, c, b, d, f, t)$, $C = (s, e, c, a, t)$, and $D = (s, a, d, f, t)$. These paths are 0.5-LO, 1.5-UBS, and bypass more than half of Opt. Path B is slightly longer than C and D, but shares with Opt only end points. For some choices of f, B is the best alternative path, and it is easy to come up with scenarios when the user will prefer B. All three paths are reasonable alternatives.

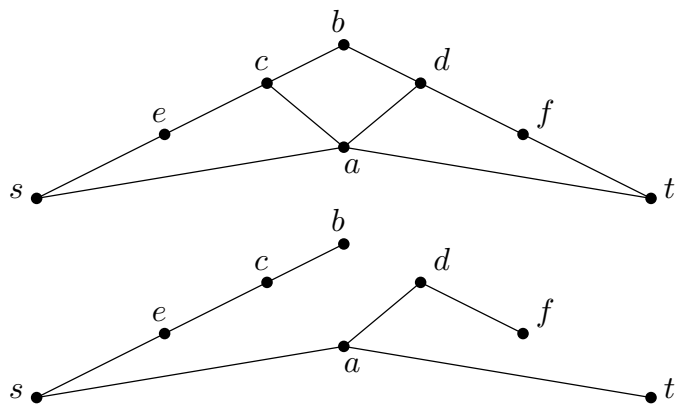

Fig. 5. Example graph with Euclidean arc lengths (top) and the shortest path tree from s (bottom). (The tree from t is its mirror image.)

In this example, however, the only non-trivial plateaus are (e, c) and (d, f). They are relatively small, and depending on parameter settings CR may not choose them. Path B corresponds to a trivial (zero-length) plateau and CR will never choose it. Thus BDV may find reasonable alternative paths even when CR fails.

A.1. Further Example Queries

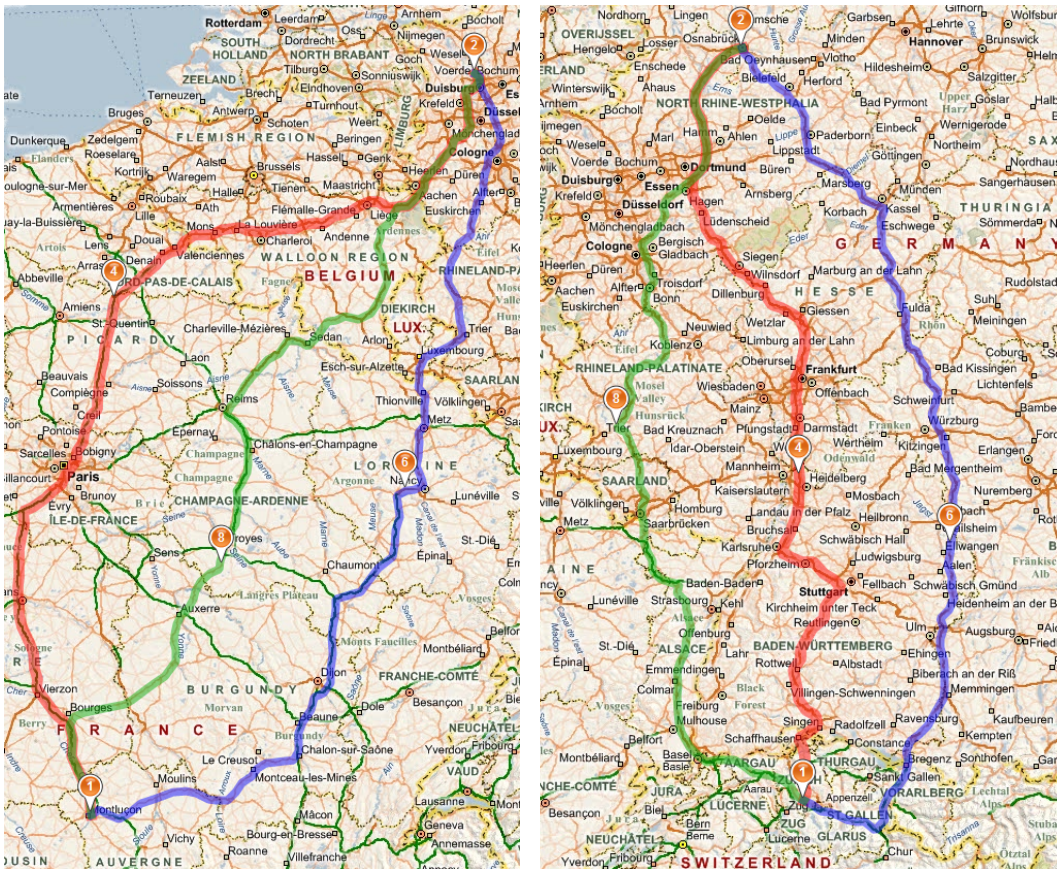
Rotterdam
SOUTH HOLLAND
NORTH BRABANT
ZEELAND
FLEMISH REGION
WALLOON REGION
BELGIUM
LUX.
DIEKIRCH
Cologne
Paris
ÎLE-DE-FRANCE
CHAMPAGNE-ARDENNE
PICARDY
LORRAINE
SAARLAND
BURGUNDY
FRANCE
FRANCHE-COMTÉ
AUVERGNE
NORTH RHINE-WESTPHALIA
Dortmund
Essen
Düsseldorf
Duisburg
Hannover
Frankfurt
Stuttgart
RHINELAND-PALATINATE
HESSE
GERMANY
THURINGIA
BADEN-WÜRTTEMBERG
ALSACE
SWITZERLAND
THURGAU
VORARLBERG
ST. GALLEN
LUCERNE
ZUG
GLARUS
JURA
NEUCHÂTEL
VAUD
Montluçon
Dijon
Chalon-sur-Saône
Mannheim
Heidelberg
Karlsruhe
Baden-Baden
Ulm
Augsburg
Ravensburg
Constance
Bregenz
Sankt Gallen
Chur
Lucerne
Zurich
Basel
Strasbourg
Metz
Nancy
Luxembourg
Trier
Koblenz
Bonn
Cologne
Aachen
Liège
Brussels
Antwerp
Amsterdam
Utrecht
Arnhem
Nijmegen
Eindhoven
Tilburg
Breda
Dordrecht
Leeuwarden
Groningen
Osnabrück
Bielefeld
Paderborn
Kassel
Eschwege
Fulda
Würzburg
Nuremberg
Erlangen
Augsburg
Ulm
Stuttgart
Heidelberg
Mannheim
Frankfurt
Wiesbaden
Mainz
Darmstadt
Kaiserslautern
Saarbrücken